732A73: Bayesian Learning

Computer Lab 3

Oriol Garrobé, Dávid Hrabovszki

08 May 2020

Question 1. Normal model, mixture of normal model with semi-conjugate prior

The data rainfall.dat consist of daily records, from the beginning of 1948 to the end of 1983, of precipitation (rain or snow in units of 1 100 inch, and records of zero precipitation are excluded) at Snoqualmie Falls, Washington.

(a)

First we analyze the date using a normal model. We assume that the daily precipitation $y_1, ..., y_n$ values are independent normally distributed:

$$y_1, ..., y_n | \mu, \sigma^2 \sim \mathcal{N}(\mu, \sigma^2)$$

where μ, σ^2 are uknown. The prior distributions for these parameters are:

$$\mu \sim N(\mu_0, \tau_0^2)$$

$$\sigma^2 \sim Inv - \chi^2(\nu_0, \sigma_0^2)$$

(i)

Now we implement a Gibbs sampler that simulates from the joint posterior distribution:

$$p(\mu, \sigma^2 | y_1, ..., y_n)$$

.

The full conditional posteriors:

$$\mu | \sigma^2, x \sim \mathcal{N}(\mu_n, \tau_n^2)$$

$$\sigma^2 | \mu, x \sim Inv - \chi^2 \left(\nu_n, \frac{\nu_0 \sigma_0^2 + \sum_{i=1}^{n} (x_i - \mu)^2}{n + \nu_0} \right)$$

where

$$\mu_n = w\bar{x} + (1 - w)\mu_0$$

,

$$w = \frac{\frac{n}{\sigma^2}}{\frac{n}{\sigma^2} + \frac{1}{\tau_0^2}}$$

,

$$\frac{1}{\tau_n^2} = \frac{n}{\sigma^2} + \frac{1}{\tau_0^2}$$

We choose the following prior values: $\mu_0 = 30$, from looking at the data, because we had no prior knowledge in the topic. To express our uncertainty about this, we set the variance of the prior mean to a high value: $\tau_0^2 = 100$. Similarly, we choose $\sigma_0^2 = 1547.103$, which is the variance of the data sample. But then we pick a small prior for the degrees of freedom, because the prior is non-informative: $\nu_0 = 4$.

(ii)

In this part we analyze the daily precipitation using our Gibbs sampler.

Figure 1 shows the trajectory of the Markov chain for sampling the mean and the cumulative mean of the mean statistic as the number of iterations grows. We ran the sampler for 1000 iterations, but it seems clear that it converges after a very short time. We ignore the first 10 draws, as this seems to be a burn-in period, where the draws are not inside the distribution. The observed mean from the data is 32.283, while the mean of the draws is 32.285, so this indicates a successful Gibbs sampling. The Markov chain has a highly oscillating shape and the fact that the cumulative mean also converges to 32.28 means that the chain converges. The daily precipitation then has an average of around 32.28, but there is a possibility that it is between 31 and 33.5.
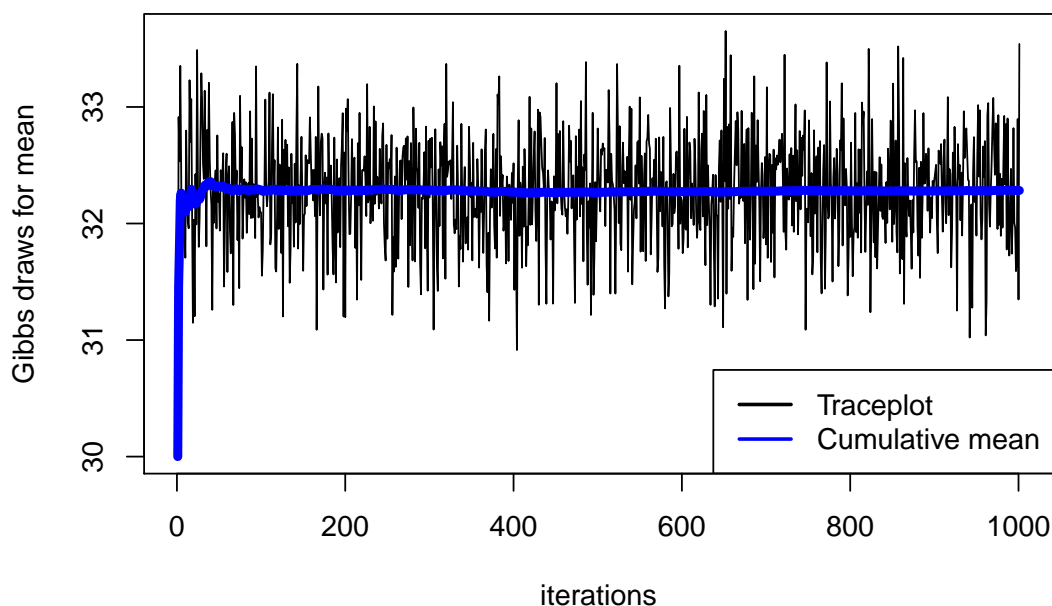

Figure 1: Markov chain of the sampled mean

Figure 2 shows the trajectory of the Markov chain for sampling the variance and the cumulative mean of the variance statistic as the number of iterations grows. We observe that the chain converges a bit more slowly than in the case of the mean, but it does so quite quickly nevertheless. Again, we ignore the burn in period (first 10 draws). The observed variance from the data is 1547.103, while the mean of the draws is 1546.152, so this indicates a successful Gibbs sampling. The Markov chain has a highly oscillating shape and the fact

2

that the cumulative mean of the draws also converges to 1546.149 means that the chain converges. The daily precipitation then has a variance of around 1546.152, but there is a possibility that it is between 1475 and 1625.
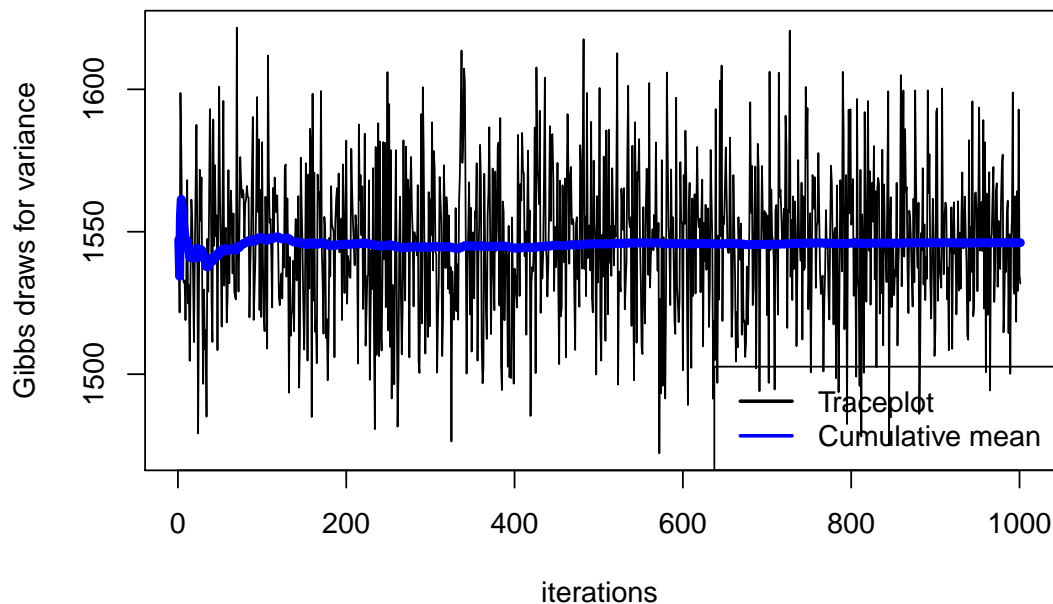

Figure 2:   Markov chain of the sampled variance

Figures 3 and 4 show that the Gibbs samples for the mean and variance follow normal distributions, which is no surpise, since the priors were normal.
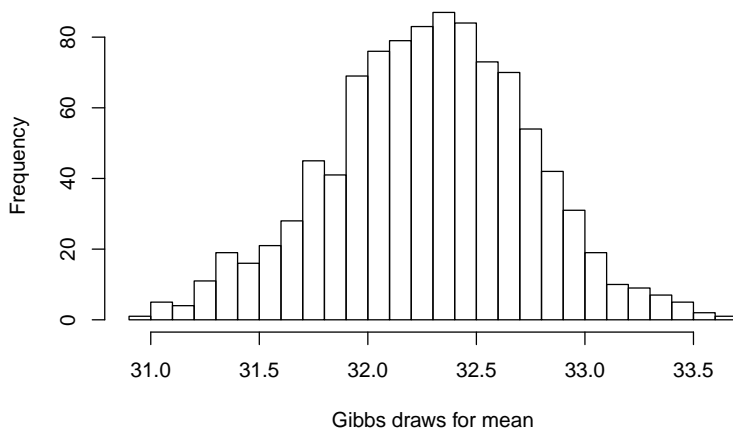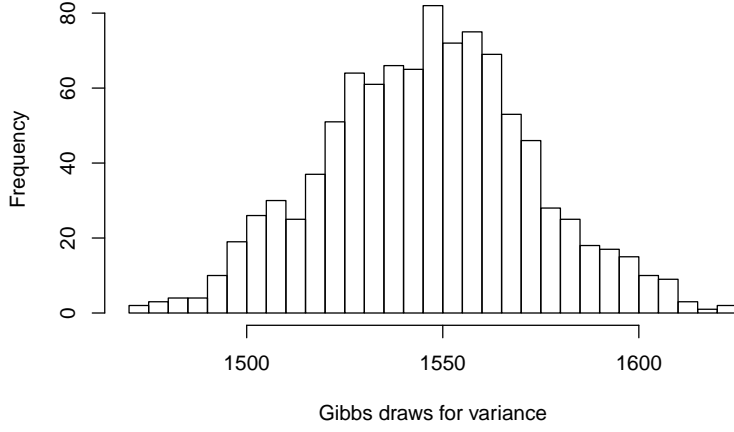

Figure 3:   Histogram of the sampled mean

Figure 4: Histogram of the sampled variance

(b)

Here we assume that the daily precipitation $y_1, ..., y_n$ follows an iid. 2-component muxture of normals model:

$$p(y_i|\mu, \sigma^2, \pi) = \pi\mathcal{N}(y_i|\mu_1, \sigma_1^2) + (1 - \pi)\mathcal{N}(y_i|\mu_2, \sigma_2^2)$$
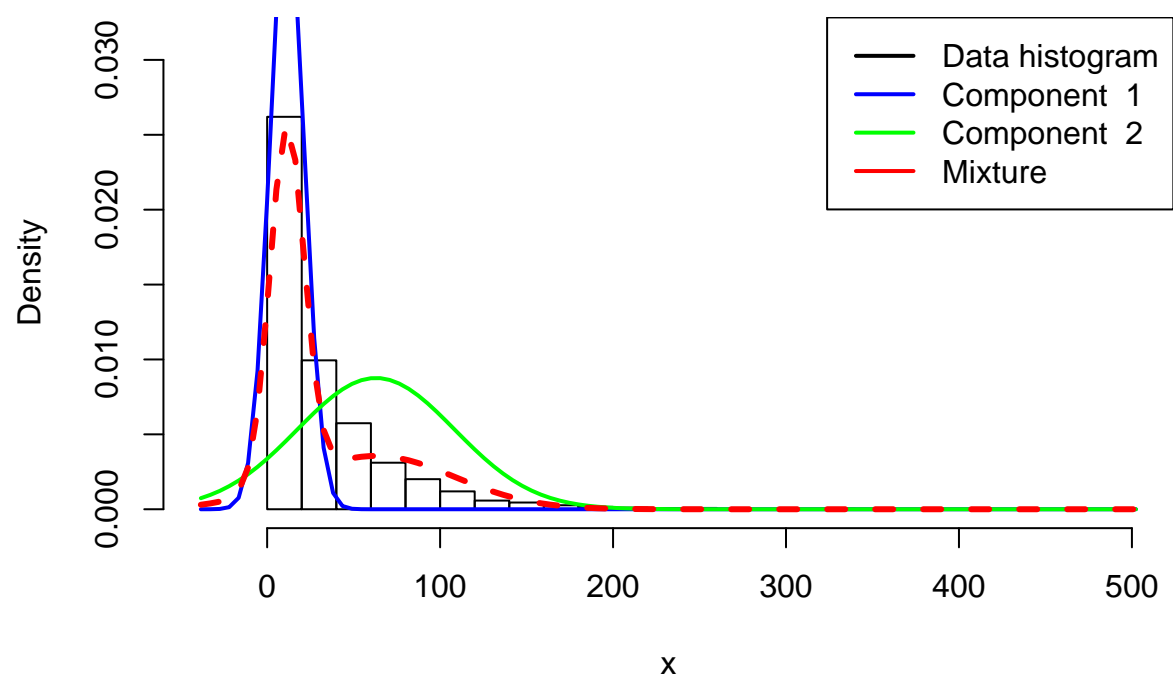
where $\mu = (\mu_1, \mu_2)$ and $\sigma^2 = (\sigma_1^2, \sigma_2^2)$

We use a slightly modified code from Mattias Villani from the file NormalMixtureGibbs.R to draw Gibbs samples from the mixture model. We set parameters the same way as in part a).
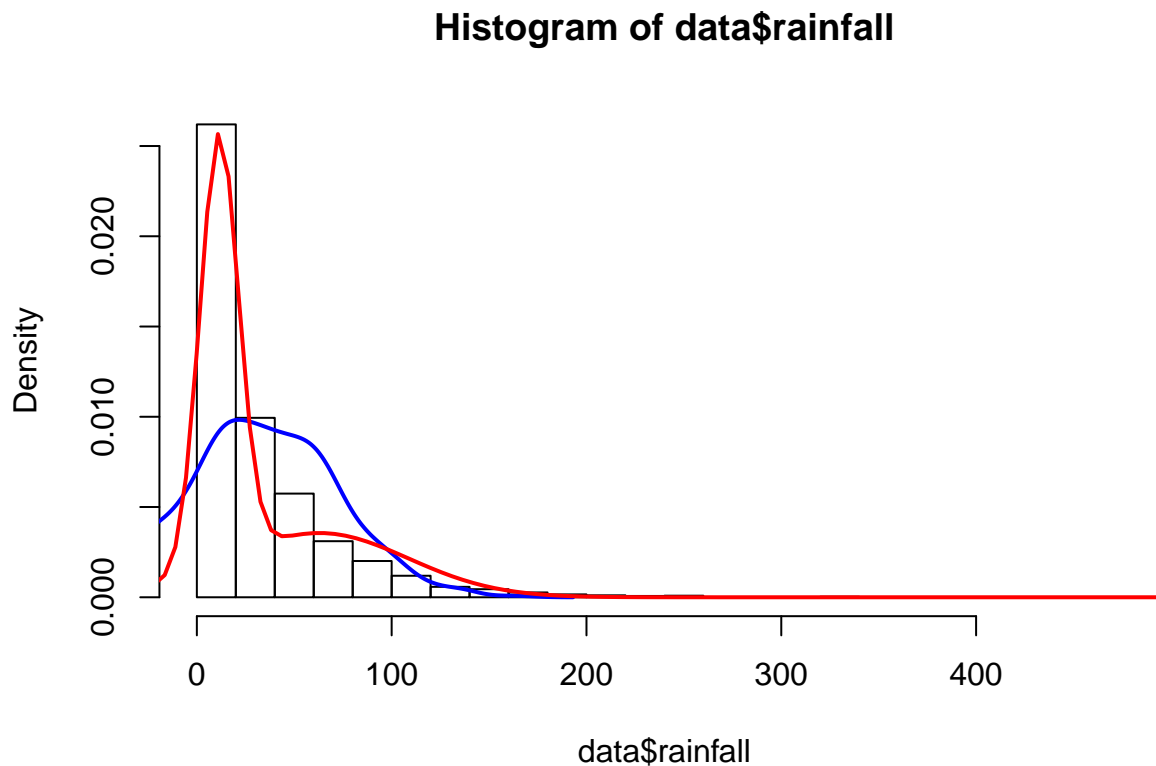
4

Histogram of x

Iteration number 100

(c)

Histogram of data$rainfall
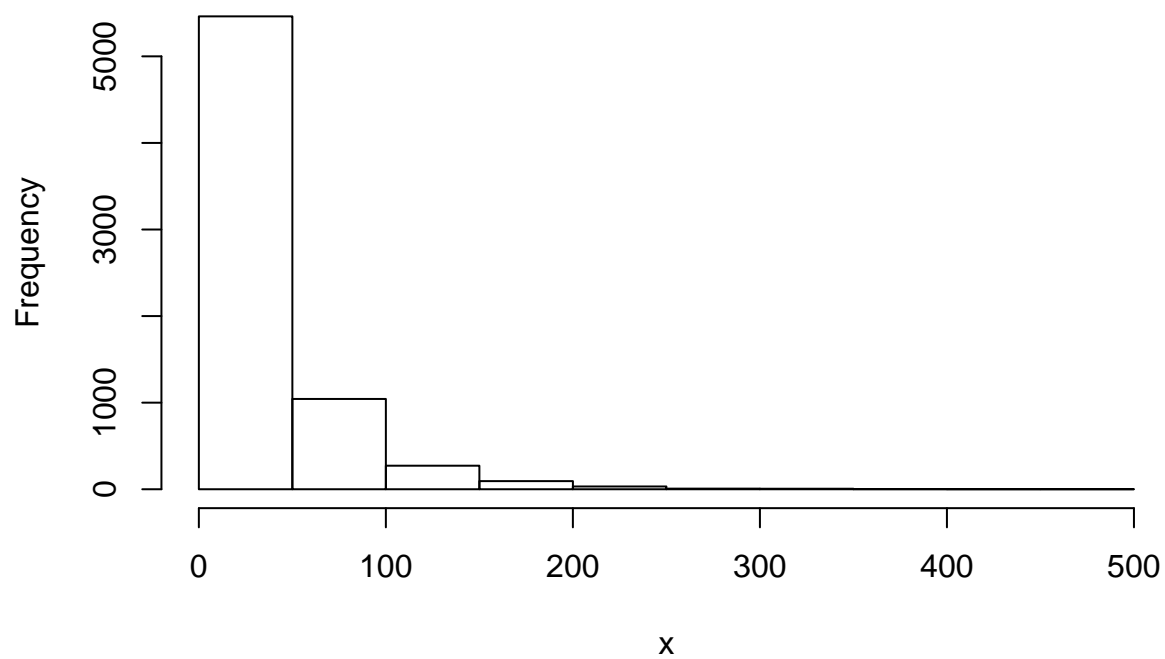
Question 2. Metropolis Random Walk for Poisson regression

In this task we consider the following Poisson regression model to analyze the dataset about coin auctions on eBay:

$$y_i|\beta \sim Poisson[exp(x_i^T\beta)], i = 1, ..., n$$

The response variable (y) is the number of bids and there are 8 features and a constant for the intercept (x). The dataset consists of 1000 auctions.

(a)

To obtain the maximum likelihood estimator of β, we use the `glm()` function.

```
##
## Call:
## glm(formula = nBids ~ . - Const, family = "poisson", data = data2)
##
## Deviance Residuals:
##     Min       1Q   Median       3Q      Max
## -3.5800  -0.7222  -0.0441   0.5269   2.4605
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
```

```
## (Intercept)   1.07244    0.03077  34.848  < 2e-16 ***
## PowerSeller -0.02054    0.03678  -0.558   0.5765
## VerifyID    -0.39452    0.09243  -4.268 1.97e-05 ***
## Sealed       0.44384    0.05056   8.778  < 2e-16 ***
## Minblem     -0.05220    0.06020  -0.867   0.3859
## MajBlem     -0.22087    0.09144  -2.416   0.0157 *
## LargNeg      0.07067    0.05633   1.255   0.2096
## LogBook     -0.12068    0.02896  -4.166 3.09e-05 ***
## MinBidShare -1.89410    0.07124 -26.588  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##     Null deviance: 2151.28  on 999  degrees of freedom
## Residual deviance:  867.47  on 991  degrees of freedom
## AIC: 3610.3
##
## Number of Fisher Scoring iterations: 5
```

From the p-values in the summary of the model, we can observe that VerifyID, Sealed, MajBlem, LogBook and MinBidShare covariates are significant at 5% level. We can also see that MinBidShare has the largest influence on the number of bids a coin receives, and that this relationship is negative, meaning that if the minimum selling price is lower than the book value, more people will bid on the coin. The largest positive influence on the number of bids comes from whether the coin is sealed in an unopened envelope (Sealed) or not. Interestingly, the fact that the seller is verified by eBay has a negative effect on the number of bids.

(b)

Now we are going to do a Bayesian analysis of the Poisson regression. The prior for β is Zellner's g-prior:

$$\mathcal{N}[0, 100 * (X^T X)^{-1}]$$

and we assume that the posterior density is approximately multivariate normal:

$$\beta | y \sim \mathcal{N}(\hat{\beta}, J_y^{-1}(\hat{\beta}))$$

where $\hat{\beta}$ is the posterior mode and $J_y(\hat{\beta})$ is the negative Hessian at the posterior mode, which we obtain from the `optim()` function. We use a modified version of our code from Lab 2, which used Mattias Villani's code (https://github.com/mattiasvillani/BayesLearnCourse/raw/master/Code/MainOptimizeSpam.zip) as a template.

The probability density function of the Poisson regression model is the following:

$$p(y_i | \beta, x_i) = \frac{exp(x_i^T \beta)^{y_i} * exp(-exp(x_i^T \beta))}{y_i!}$$

Therefore the log-likelihood is:

$$logp(y | \beta, x) = \sum_{i=1}^{n} \left(y_i x_i \beta - exp(x_i \beta) - log(y_i!) \right)$$

We obtain the following posterior mode (or mean, since it's normal distribution) for β, and we can observe that it is very similar to the maximum likelihood estimator:

Const	PowerSeller	VerifyID	Sealed	Minblem	MajBlem	LargNeg	LogBook	MinBidShare
1.0698	-0.0205	-0.3932	0.4436	-0.0524	-0.2213	0.0707	-0.1202	-1.8919

(c)

In this part we use the Metropolis Random Walk algorithm to simulate from the actual posterior of β. Our implementation can generate draws from any posterior density with a vector of model parameters θ. In this case, the proposal density is multivariate normal:

$$\theta_p|\theta^{(i-1)} \sim \mathcal{N}(\theta^{(i-1)}, c * \Sigma)$$

where $\Sigma = J_y^{-1}(\hat{\beta})$ obtained in part b). We found that $c = 0.65$ provides a covariance matrix that results in an average acceptance probability of proposals between 25% and 30%.
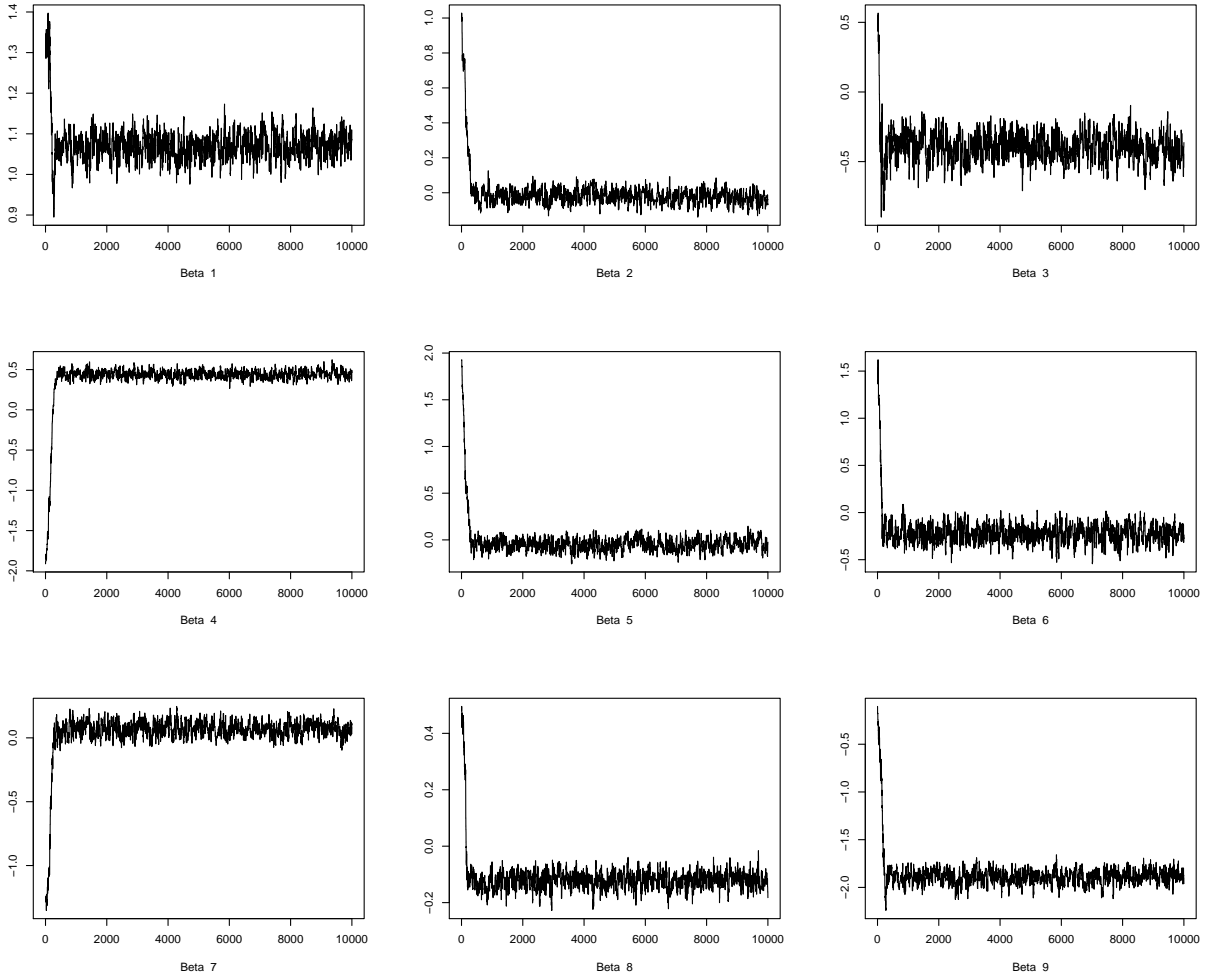

Figure 5: Markov chains of the posterior draws of the parameters

Figure 5 shows the Markov chains obtained for all β parameters. We ran the Metropolis Random Walk algorithm for 10000 iterations to make sure that the chains converge, because the burn-in period is quite

large. The convergence is satisfactory based on the plots, so we save the chains after discarding the first 1000 draws of the burn-in period. The average acceptance probability is 27%.

Calculating the means of the parameters, we obtain the following table:

Const	PowerSeller	VerifyID	Sealed	Minblem	MajBlem	LargNeg	LogBook	MinBidShare
1.0698	-0.021	-0.3944	0.4414	-0.0487	-0.2265	0.0717	-0.1216	-1.8882

This, again is very similar to the appriximate results from 2.b) and the maximum likelihood estimator from 2.a).

(d)

Now we use the MCMC draws from 2.c) to simulate from the predictive distribution of the number of bidders in a new auction with the following characteristics:

- PowerSeller = 1
- VerifyID = 1
- Sealed = 1
- MinBlem = 0
- MajBlem = 0
- LargNeg = 0
- LogBook = 1
- MinBidShare = 0.5

The probability that nobody bids on the coin in this new auction is 35.48%, which we get by dividing the number of predictive draws where the result is 0 by the number of all the draws. Or we can look at the histogram of the normalized predictive distribution in Figure 6. The low number of bids is probably because the minimum selling price for the coin was set high compared to its book value (MinBidShare = 0.5).
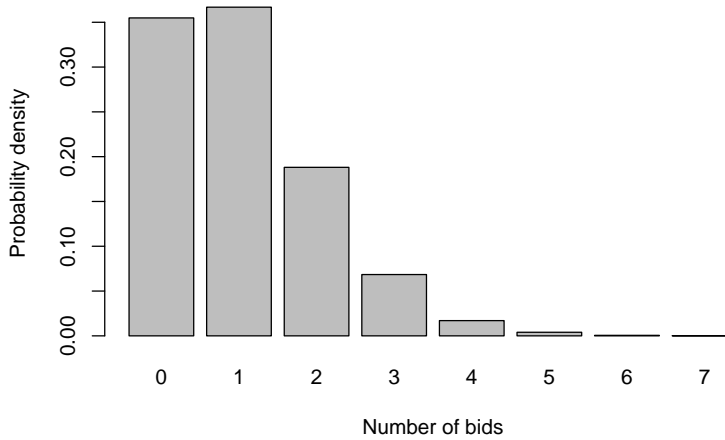

Figure 6: Histogram of the predicted distribution

Appendix

```r
knitr::opts_chunk$set(echo=FALSE, eval=TRUE)
# R version
RNGversion('3.5.1')

# libraries
library(knitr)
library(mvtnorm)

#### 1.a
data = read.table('rainfall.dat', col.names = 'rainfall')

#### (i)

n = length(data$rainfall)
avg = mean(data$rainfall)

#prior setup
mu_0 = 30 #mu prior from prior knowledge, or looking at the data if there is no knowledge
tau2_0 = 100 #var(mu_0) high, because we have no prior knowledge
sigma2_0 = var(data$rainfall) # should use var(data$rainfall for best guess)
nu_0 = 4 #df, small, because we have no prior knowledge

# scaled-inverse chi-square simulator (from lab1 and from NormalMixtureGibbs.R)
rScaledInvChi2 <- function(n, df, scale){
  return((df*scale)/rchisq(n,df=df))
}

# Gibbs sampling
steps = 1000
gibbsDraws <- matrix(0,steps+1,2) #1st column: mu, 2nd column: sigma2
gibbsDraws[1,1] = mu_0
gibbsDraws[1,2] = sigma2_0

# Seed
set.seed(1234567890)
for (i in 1:steps) {
  #simulate mu from conditional posterior (normal)
  w = (n/gibbsDraws[i,2])/(n/gibbsDraws[i,2] + 1/tau2_0)
  mu_n = w * avg + (1-w) * mu_0
  tau2_n = 1 / (n/gibbsDraws[i,2] + 1/tau2_0)

  #update mu
  gibbsDraws[i+1,1] = rnorm(1, mu_n, sqrt(tau2_n))

  #simulate sigma2 from conditional posterior (inv-chisq)
  nu_n = nu_0 + n
  scale = (nu_0 * sigma2_0 + sum((data$rainfall - gibbsDraws[i+1,1])^2)) / (n + nu_0)

  #update sigma2
  gibbsDraws[i+1,2] = rScaledInvChi2(1, nu_n, scale)
}
```

```r
#### (ii)
#Plot Markov chains and cumulative mean plots
cum_mean = numeric(steps + 1)
cum_var = numeric(steps + 1)
for (i in 1:(steps+1)) {
  cum_mean[i] = mean(gibbsDraws[1:i,1])
  cum_var[i] = mean(gibbsDraws[1:i,2])
}
plot(1:(steps+1), gibbsDraws[,1], type = 'l', xlab = 'iterations', ylab = 'Gibbs draws for mean')
legend(x='bottomright', legend=c('Traceplot', 'Cumulative mean'), col = c('black', 'blue'), lwd = 2)
lines(1:(steps+1), cum_mean, col = 'blue', lwd = 5)
plot(1:(steps+1), gibbsDraws[,2], type = 'l', xlab = 'iterations', ylab = 'Gibbs draws for variance')
legend(x='bottomright', legend=c('Traceplot', 'Cumulative mean'), col = c('black', 'blue'), lwd = 2)
lines(1:(steps+1), cum_var, col = 'blue', lwd = 5)

# result: mean of all posterior draws
burn_in_mu = 10
burn_in_sigma2 = 10
mu_gibbs = mean(gibbsDraws[burn_in_mu:length(gibbsDraws[,1]),1]) #32.28486
sigma2_gibbs = mean(gibbsDraws[burn_in_sigma2:length(gibbsDraws[,2]),2]) #1546.152

hist(gibbsDraws[burn_in_mu:length(gibbsDraws[,1]),1], breaks = 30, main = '', xlab = 'Gibbs draws for m

hist(gibbsDraws[burn_in_sigma2:length(gibbsDraws[,2]),2], breaks = 40, main = '', xlab = 'Gibbs draws f

#### 1.b

#mixture model (normals)
# Code from NormalMixtureGibbs.R (by Mattias Villani)

# Estimating a simple mixture of normals
# Author: Mattias Villani, IDA, Linkoping University. http://mattiasvillani.com

##########    BEGIN USER INPUT ################
# Data options
x <- as.matrix(data$rainfall)

# Model options
nComp <- 2    # Number of mixture components

# Prior options
alpha <- 10*rep(1,nComp) # Dirichlet(alpha)
muPrior <- rep(30,nComp) # Prior mean of mu
tau2Prior <- rep(100,nComp) # Prior std of mu: high, because we have no prior knowledge
sigma2_0 <- rep(var(x),nComp) # s20 (best guess of sigma2)
nu0 <- rep(4,nComp) # degrees of freedom for prior on sigma2: small, because we have no prior knowledge

# MCMC options
nIter <- 100 # Number of Gibbs sampling draws

# Plotting options
plotFit <- TRUE
lineColors <- c("blue", "green", "magenta", 'yellow')
```

```
sleepTime <- 0.1 # Adding sleep time between iterations for plotting
################   END USER INPUT ###############

###### Defining a function that simulates from the
rScaledInvChi2 <- function(n, df, scale){
  return((df*scale)/rchisq(n,df=df))
}


####### Defining a function that simulates from a Dirichlet distribution
rDirichlet <- function(param){
  nCat <- length(param)
  piDraws <- matrix(NA,nCat,1)
  for (j in 1:nCat){
    piDraws[j] <- rgamma(1,param[j],1)
  }
  piDraws = piDraws/sum(piDraws) # Diving every column of piDraws by the sum of the elements in that co
  return(piDraws)
}


# Simple function that converts between two different representations of the mixture allocation
S2alloc <- function(S){
  n <- dim(S)[1]
  alloc <- rep(0,n)
  for (i in 1:n){
    alloc[i] <- which(S[i,] == 1)
  }
  return(alloc)
}
# Seed
set.seed(1234567890)

# Initial value for the MCMC
nObs <- length(x)
S <- t(rmultinom(nObs, size = 1 , prob = rep(1/nComp,nComp))) # nObs-by-nComp matrix with component all
mu <- quantile(x, probs = seq(0,1,length = nComp))
sigma2 <- rep(var(x),nComp)
probObsInComp <- rep(NA, nComp)

# Setting up the plot
xGrid <- seq(min(x)-1*apply(x,2,sd),max(x)+1*apply(x,2,sd),length = 100)
xGridMin <- min(xGrid)
xGridMax <- max(xGrid)
mixDensMean <- rep(0,length(xGrid))
effIterCount <- 0
ylim <- c(0,2*max(hist(x)$density))


for (k in 1:nIter){
  #message(paste('Iteration number:',k))
  alloc <- S2alloc(S) # Just a function that converts between different representations of the group al
  nAlloc <- colSums(S)
  #print(nAlloc)
  # Update components probabilities
```

```r
    pi <- rDirichlet(alpha + nAlloc)

    # Update mu's
    for (j in 1:nComp){
      precPrior <- 1/tau2Prior[j]
      precData <- nAlloc[j]/sigma2[j]
      precPost <- precPrior + precData
      wPrior <- precPrior/precPost
      muPost <- wPrior*muPrior + (1-wPrior)*mean(x[alloc == j])
      tau2Post <- 1/precPost
      mu[j] <- rnorm(1, mean = muPost, sd = sqrt(tau2Post))
    }

    # Update sigma2's
    for (j in 1:nComp){
      sigma2[j] <- rScaledInvChi2(1, df = nu0[j] + nAlloc[j], scale = (nu0[j]*sigma2_0[j] + sum((x[alloc =
    }

    # Update allocation
    for (i in 1:nObs){
      for (j in 1:nComp){
        probObsInComp[j] <- pi[j]*dnorm(x[i], mean = mu[j], sd = sqrt(sigma2[j]))
      }
      S[i,] <- t(rmultinom(1, size = 1 , prob = probObsInComp/sum(probObsInComp)))
    }

    # Printing the fitted density against data histogram
    if (k == nIter) {
      if (plotFit && (k%%1 ==0)){
      effIterCount <- effIterCount + 1
      hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax), main = paste("Iteration number",k),
      mixDens <- rep(0,length(xGrid))
      components <- c()
      for (j in 1:nComp){
        compDens <- dnorm(xGrid,mu[j],sd = sqrt(sigma2[j]))
        mixDens <- mixDens + pi[j]*compDens
        lines(xGrid, compDens, type = "l", lwd = 2, col = lineColors[j])
        components[j] <- paste("Component ",j)
      }
      mixDensMean <- ((effIterCount-1)*mixDensMean + mixDens)/effIterCount

      lines(xGrid, mixDens, type = "l", lty = 2, lwd = 3, col = 'red')
      legend("topright", box.lty = 1, legend = c("Data histogram",components, 'Mixture'),
      col = c("black",lineColors[1:nComp], 'red'), lwd = 2)
      #Sys.sleep(sleepTime)
    }
    }


}

# hist(x, breaks = 20, freq = FALSE, xlim = c(xGridMin,xGridMax), main = "Final fitted density")
# lines(xGrid, mixDensMean, type = "l", lwd = 2, lty = 4, col = "red")
```

```r
# lines(xGrid, dnorm(xGrid, mean = mean(x), sd = apply(x,2,sd)), type = "l", lwd = 2, col = "blue")
# legend("topright", box.lty = 1, legend = c("Data histogram","Mixture density","Normal density"), col=

########################    End of code from Mattias Villani    ####################################
#### c

#histogram of data
hist(data$rainfall, freq = F, breaks = 30)

#normal density from a)
norm_gibbs = rnorm(1000, mean = mu_gibbs, sd = sqrt(sigma2_gibbs))
lines(density(norm_gibbs), col = 'blue', lwd = 2)

#mixture of normals from b)
lines(xGrid, mixDensMean, type = "l", lwd = 2, col = "red")
################### QUESTION 2 ####################

data2 = read.table('eBayNumberOfBidderData.dat', header = T)


#2.a
model = glm(nBids ~ . - Const, data = data2, family = 'poisson')
summary(model)

#2.b
# The following code was written using Mattias Villani's implementation as a template:
# https://github.com/mattiasvillani/BayesLearnCourse/raw/master/Code/MainOptimizeSpam.zip

y = as.vector(data2[,1])
x = as.matrix(data2[,2:10])
nfeatures = dim(x)[2]

# prior
mu = as.vector(rep(0,nfeatures))
sigma = 100*solve(t(x)%*%x)

LogPrior = function(theta, mu, sigma){
  return(dmvnorm(theta, as.matrix(mu), sigma, log = T))
}

# log likelihood
LogLikelihood = function(theta, y, x){
  sum(y * x%*%theta - exp(x%*%theta) - log(factorial(y))) #log-likelihood of poisson regression
}

# log posterior
LogPosterior = function(theta, y, x, mu, sigma){
  LogPrior(theta, mu, sigma) + LogLikelihood(theta, y, x)
  # we can sum them instead of multiplying, because of the log
}

# initialize Beta vector randomly
# Seed
```

```r
set.seed(1234567890)
Beta_init = as.vector(rnorm(dim(x)[2]))

# optimizing the log posterior by changing the Betas (maximize)
res = optim(Beta_init, LogPosterior, gr = NULL, y, x, mu, sigma,
            method="BFGS", control=list(fnscale=-1), hessian=T)

Beta_hat = res$par # posterior mode
Hessian = res$hessian
post_sigma = solve(-Hessian) # posterior cov matrix
stdev = sqrt(diag(post_sigma))

### end of code from Mattias Villani's template

betas = t(as.matrix(round(Beta_hat,4)))
colnames(betas) = colnames(data2[2:10])
kable(betas)
#2.c

RWMSampler = function(logPostFunc, c, iter, initial, proposal_sigma, ...){
  X = matrix(0,nrow = iter+1, ncol = length(initial)) #returned posterior draws
  X[1,] = initial

  #acceptance probability stats
  alphas = numeric(iter)

  for (t in 1:iter) {
    #print(t)
    Y = as.numeric(rmvnorm(1, mean = X[t,], sigma = c * proposal_sigma)) # Proposal needs to be a numer
    U = runif(1, min = 0, max = 1) # Generate uniform

    alpha = min(1,
                (exp(logPostFunc(Y, ...) - logPostFunc(X[t,], ...)))) #have to take exp(), because post

    if (U < alpha) {
      X[t+1,] = Y #accept new proposal with probability alpha
    } else {
      X[t+1,] = X[t,] #reject new proposal
    }
    alphas[t] = alpha # save acceptance probability
    t = t+1
  }
  return(list('mcmc' = X, 'alphas' = alphas))
}

# Seed
set.seed(1234567890)
nIter = 10000
mc = RWMSampler(LogPosterior, c = 0.65, iter = nIter, initial = Beta_init, proposal_sigma = post_sigma,
par(mfrow=c(3,3))
for (p in 1:9) {
  plot(mc$mcmc[,p], type = 'l', ylab = '', xlab = paste('Beta ', p)) #plot mcmc-s for all betas
}
```

```r
burn_in_theta = 1000
#mean(mc$alphas) #0.27 average acceptance probability

Beta_mcmc = mc$mcmc[burn_in_theta:nIter,] # discard burn-in period
betas_mcmc = t(as.matrix(round(colMeans(Beta_mcmc),4)))
colnames(betas_mcmc) = colnames(data2[2:10])
kable(betas_mcmc)


#2.d
Const = 1
PowerSeller = 1
VerifyID = 1
Sealed = 1
MinBlem = 0
MajBlem = 0
LargNeg = 0
LogBook = 1
MinBidShare = 0.5
x_pred = c(Const, PowerSeller, VerifyID, Sealed, MinBlem, MajBlem, LargNeg, LogBook, MinBidShare)

predictive_dist = numeric(dim(Beta_mcmc)[1])
# Seed
set.seed(1234567890)
for (i in 1:dim(Beta_mcmc)[1]) {
  predictive_dist[i] = rpois(1, exp(t(as.matrix(x_pred)) %*% as.matrix(Beta_mcmc[i,]))) #sample poisson
}

#sum(predictive_dist==0)/length(predictive_dist) #0.3548495

#normalized predictive distribution
values = unique(predictive_dist)
values = cbind(values,numeric(length(values)))
for (v in values[,1]) {
  count = sum(predictive_dist==v)
  values[values[,1]==v] = c(v, count)
}

values = data.frame(values)
colnames(values) = c('value', 'count')
values$ratio = values$count/sum(values$count)
values = values[order(values$value),]

barplot(values$ratio, names.arg = values$value, xlab = 'Number of bids', ylab = 'Probability density')
```